



Creditworthy SMEs get rejected for fixable reasons.

Product and technology overview — for a software developer and an investor

A credit-readiness advisory for German and Austrian SMEs (10–249 employees). It diagnoses exactly which factors pull a company's file below a lender's approval line, separates presentation problems from genuine credit risk, and routes the corrected file to the lender type most likely to approve it. A free AI-assisted check brings companies in; paid analysis and advisory follow.

20 Sept 2026

Version 0.3.0

Working prototype ~9,000 lines

No paying clients yet

Python, 217 tests passing

Germany first

Austria second

00 How to read this

If you are...	Read
An investor	\$1 Problem · \$2 Solution · \$3 How it works · \$4 AI · \$7 Challenges · \$8 Costs · \$9 Goals · \$10 Status
A software developer	\$3 How it works · \$5 Stack and architecture · \$6 What is left to build · \$7.1 Technical challenges · \$11 Running it locally

Everything in this document is stated at its real confidence level. Figures sourced from public data name their source. Anything marked **ESTIMATE** is a planning assumption, not a measurement. The product has been validated on fictional cases only — that is the single most important caveat on every page that follows.

01 The problem

German SME lending is tightening structurally, and the companies caught by it have nowhere to turn between their bookkeeper and a consultancy priced for corporates.

40.5% report restrictive banks

Share of Mittelstand firms in credit negotiations reporting restrictive bank behaviour — a record since the survey began in 2017, up 6.5 points in a single quarter. *KfW Research, 7 July 2026.*

Germany-specific worsening

SME rejection rates rose further in Germany through Q3–Q4 2025, faster for SMEs than for large firms, while other large euro-area countries were flat or falling. *ECB Bank Lending Survey.*

Nobody serves the middle

Bookkeepers do not read rating models. Premium consultancies are priced for €50M+ companies. Automated fintech lenders do not advise. The €3–15M SME is unserved.

What actually happens to the company

- The SME learns it has a problem **only after the bank has said no** — and the rejection reason it is given is usually a sentence, not a diagnosis.
- A logged rejection makes the next application harder. Reapplying without changing anything is the most common and most expensive mistake.
- Much of what fails is **not the business**: a shareholder loan without a subordination agreement reads as debt; an overdraft used as permanent financing reads as distress; a management report six months old reads as weak reporting. Each is fixable in weeks.
- Nobody currently tells the owner which of the two situations they are in.

02 The solution

Tell the company, before it applies, exactly which factors fail and whether they are fixable — then either fix them or say honestly that they cannot be fixed.

Three products, one funnel

Tier	Price	What the company gets	Role in the business
Credit check	Free	Uploads its annual accounts. The figures are read automatically, the company confirms each one and answers ~12 questions. Instant readiness band (A–E), its three biggest weaknesses, and a plain-language explanation.	Acquisition and qualification. Self-serve, near-zero marginal cost.
Full report	€390	Full questionnaire (41 questions) plus the tax advisor's 20, the document set, the complete diagnostic, the before/after simulation and the lender-type routing — reviewed and released by a human advisor.	Mostly automated, low marginal cost. The volume product.
Advisory	€1,500–3,000	Calls, coordination with the company's tax advisor, sign-off on reclassifications, support up to the bank meeting.	Flat fee per engagement, earned regardless of outcome.
Success fee OPTIONAL	1.5–2%	Only where placement is actively supported.	Likely requires a §34c GewO licence — a decision deliberately not yet taken.

All fee figures are planning assumptions. The free check widens the addressable funnel from the ~16,000 companies that hit obstacles each year to the ~113,000 that apply (§9).

Two functions, deliberately kept separate

1 · Diagnostic engine

A transparent rules-and-ratios engine — DSCR, equity ratio, leverage, margins, payment behaviour — flags which factors pull the file below a typical approval line and simulates the effect of each fix. Sold standalone at a flat fee.

2 · Remediation and routing

An advisor turns each flag into an action — reclassify a shareholder loan, fix reporting cadence, swap an overdraft for a term loan or factoring, use a KfW liability-release programme — and matches the corrected profile to a lender type.

The separation keeps the core business **advisory** — not brokerage, not a rating agency. The SME pays, never the lender, so there is no conflict between "best terms for the client" and "whoever pays the most referral fee".

The distinction the whole business rests on: we say no. Every weakness is classified as presentation, documentation, product fit, collateral — or genuine credit risk. Genuine risk is never simulated away; rule R08 refuses, and one substantive finding outranks any number of cosmetic ones. Those clients are advised not to apply. Polishing a distressed file produces a better-looking rejection, burns the lender relationships and poisons the outcome dataset that is meant to become the moat. Declining the engagement is the product working correctly, and the demo includes such a case deliberately.

03 How it works, end to end

The free check — four steps, a few minutes

1 · Upload Annual accounts as PDF or scan. Nothing else needed.	2 · Read 33 balance-sheet and P&L fields extracted automatically.	3 · Confirm Every figure shown for the company to check and correct. Consistency checks must pass.	4 · Result ~12 questions, then band A–E, top three weaknesses, plain- language explanation.
--	--	--	---

The paid engagement — the full pipeline

Stage	What happens	Built from
Intake	Two questionnaires — 41 questions for the company, 20 for its tax advisor; every question states <i>why</i> it is asked. A catalogue of 14 document types says who supplies each, whether it is required or conditional, and in which format.	<code>intake/questionnaire.py</code> , <code>intake/documents.py</code>
Import	Automatic reading of the two structured sources: the DATEV trial balance (current and prior year) and bank-statement CSVs, which yield days at the overdraft limit and returned direct debits.	<code>ingest/datev.py</code> , <code>ingest/bank_csv.py</code>
Assembly	One case from three parties with a strict precedence — bank data beats the tax advisor, the tax advisor beats the client — and a provenance record for every contested figure.	<code>intake/assemble.py</code>
Validation	Refuses to analyse a file whose balance sheet does not balance. A confident diagnostic from a broken input is this business's worst failure mode.	<code>validation.py</code>
Ratios	A 26-value ratio set computed the way a German bank's rating engine does it, annualised correctly when working from a mid-year management report.	<code>ratios.py</code>
Score	Transparent weighted scorecard: 10 factors, piecewise-linear over published thresholds, renormalised when a factor is unavailable. Every factor exposes its value, weight and points lost. Bands A–E.	<code>scorecard.py</code>
Diagnose	10 fixability rules (R01–R10) classify each weakness and attach the concrete measure, its cost and its lead time.	<code>remediation.py</code>
Simulate	Applies the measures to a copy of the case and recomputes: before/after score and band.	<code>remediation.py</code>
Route	Ranks 7 lender types — house bank, KfW with liability release, state guarantee bank, factoring, leasing, digital lender, mezzanine/MBG — for both the current and the corrected profile.	<code>routing.py</code>
Report	German client report with the disclaimer on every page, plus the request letters and the tax-advisor sign-off sheet. Released only after a human advisor reviews it.	<code>reporting/</code>
Outcome	One row per engagement — weaknesses flagged, measures applied, lender routed to, approved or not — appended to the outcome log.	<code>workflow.py</code>

The outcome log is the long-term asset. Early on it is a spreadsheet, not a model. It is what eventually moves the scorecard thresholds from published convention to something calibrated on this business's own placement history — a dataset no competitor can buy.

Three parties, three views

Role	Sees
Company (SME)	A guided seven-step flow: its own case, its questionnaire, its uploads, and the report once released.
Tax advisor	Only the cases of clients who invited them. Confirmations, uploads, the sign-off sheet — not the company's questionnaire answers.

Role	Sees
Advisor (us)	The whole pipeline: every case, run the diagnostic, release reports, generate letters, record outcomes.

Permissions are enforced server-side, not hidden in the interface, and are covered by tests that exercise them over real HTTP with cookies.

04 Where the AI sits — and where it deliberately does not

"AI reads and explains. Transparent rules decide." That sentence is the product's positioning, its legal posture and its architecture, and it is enforced in code.

Task	Who does it	Why
Read balance sheet and P&L from a PDF or scan	AI (Claude) or the local rules reader	Removes the real friction: typing figures. Works on any layout.
Decide which figures are used	The company , confirming each figure	A proposal is not an input until a person has confirmed it.
Check the figures	Deterministic checks	Assets = liabilities; P&L result = balance-sheet result. Catches misreads before they reach a ratio.
Band, findings, simulation, routing	Rules engine only	Explainable, auditable, legally defensible. Never a model.
Explain the result, answer questions	AI , grounded only in the engine's output	An owner understands "your equity ratio reads as 10.1% because...", not a ratio table.
Guardrail on every AI sentence	Deterministic filter	Rejects approval promises, odds and rating vocabulary; falls back to a fixed template.
Release of the paid report	A human advisor	Nothing paid-for reaches a client unreviewed.

Why this split, and not "an AI that scores companies"

- **Regulatory.** A transparent rules engine keeps the product on the advisory side of the EU Credit Rating Agencies Regulation, and outside the EU AI Act's high-risk "creditworthiness of natural persons" category. An opaque score that ranks creditworthiness is much closer to a rating.
- **Saleability.** The buyer is an SME owner and their tax advisor. A model output cannot be argued with by an accountant; a ratio can — and being argued with is how the advice earns its fee.
- **Data reality.** At two to three engagements a month there will be tens of outcomes in year one, not thousands. There is no training set. A model fitted to that would be noise wrapped in false authority.

Recorded as ADR-001 in the repository: replacing the engine with a model is a legal decision, not a technical upgrade, and must not happen as an incremental refactor.

Data protection built into the AI path

Control	Status	Detail
Local reader is the default — no data leaves the server	IN PLACE	Claude is opt-in per deployment via one environment variable
Separate, voluntary AI consent per company	IN PLACE	Without it the document is never sent; manual entry offered instead
Data minimisation towards the AI	IN PLACE	Document only for reading; for explanations only the result summary — no names, no documents
Audit log of every AI call	IN PLACE	Time, purpose, provider, model, SHA-256 of the input, actor — never the content
Self-service erasure (Art. 17 GDPR)	IN PLACE	Deletes account, cases and documents

Control	Status	Detail
Transparency labelling (EU AI Act Art. 50)	IN PLACE	AI-read figures and AI-written text are labelled as such in the interface
Data processing agreements incl. the AI provider; no training on client data	BEFORE GO-LIVE	Contractual, plus transfer basis and inference region
EU hosting, encryption at rest, DPIA, retention schedule, 2FA	BEFORE GO-LIVE	See §6 and §7.2

The full 10-point pre-launch checklist lives in `docs/ai-and-data-protection.md`, and a public-facing version of this table is already on the website at `/sicherheit`.

05 The stack

Written for a developer picking this up. The guiding decision was **zero runtime dependencies**: the entire product — engine, API and website — runs on a laptop with nothing but a Python install.

Language	Python 3.11+ (developed on 3.14)
Runtime dependencies	None. Standard library only, including the web server
Optional extras	[dev] pytest · [ai] anthropic SDK · [pdf] pymupdf for scanned documents
Web server	stdlib ThreadingHTTPServer ; a thin JSON API over the workflow layer, one file
Front end	Vanilla JavaScript SPA — no framework, no build step, no bundler. Hash routing, DOM built with <code>textContent</code> only (no <code>innerHTML</code> , so no injection surface)
Languages (UI)	German and English throughout, switchable live; German number formatting on every client-facing document
Auth	PBKDF2-SHA256, 600,000 iterations; HttpOnly / SameSite=Strict session cookies; login lockout; strict CSP; origin checks
Storage today	JSON case folders on disk behind a <code>CaseStore</code> interface — the designated database swap point
AI (optional)	Anthropic Claude (<code>claude-opus-5</code>) for document reading and explanations, behind a provider interface with a local rules-based fallback
Tests	pytest — 217 tests , all passing; includes real-HTTP portal tests, regulatory-wording tests and a calibration check against the Bundesbank dataset
Size	~9218 lines of Python across 38 modules

Module map

src/credit_readiness/	
└ models.py	HGB §266/§275 domain objects; economic-equity restatement
└ validation.py	input plausibility; errors block analysis by default
└ ratios.py	ratio computation, partial-year annualisation
└ scorecard.py	factor weights, breakpoints, banding [REGULATORY]
└ remediation.py	fixability rules R01-R10 + simulation [CORE IP]
└ routing.py	lender-type matching [REGULATORY]
└ benchmarks.py	sector medians [PLACEHOLDER DATA]
└ engine.py	orchestrator → DiagnosticResult
└ ai/	document reading (rules Claude), explanations, guardrails, audit log
└ intake/	questionnaires, document catalogue, case assembly
└ ingest/	json_intake (canonical), datev (SKR04), bank_csv
└ casefile.py	CaseStore interface + local folder store [SWAP FOR DATABASE]
└ auth.py	accounts, roles, sessions (PBKDF2) [SWAP FOR DATABASE]
└ workflow.py	overview, diagnose, letters, outcome – used by portal and CLI
└ reporting/	report, HTML, letters, printable forms, summary JSON
└ webapp/	website, investor page, portal (API + front end, DE/EN)
└ cli.py	diagnose / batch / datev / serve / case / questionnaire / bank

Why zero dependencies was a deliberate choice

- The product handles confidential financial data on a founder's machine before it is ever hosted. A small dependency tree is a small supply-chain surface.
- Everything is inspectable. There is no framework magic between an HTTP request and the rule that produced a number — which matters when a regulator or a tax advisor asks how a figure was derived.
- It is not dogma: the three places where a real deployment needs real infrastructure are named in the code (`CaseStore` , `auth.py` , `webapp/server.py`), and are additions rather than rewrites.

Everything the portal does is also a command

```
python -m credit_readiness serve --demo --open          # website + portal, fictional data
python -m credit_readiness case new "Firma GmbH"
python -m credit_readiness case upload CRA-2026-0001 susa_aktuell susa.csv
python -m credit_readiness case diagnose CRA-2026-0001
python -m credit_readiness batch data/samples           # whole folder of cases
```

Useful for the developer: the entire business logic is reachable without the web layer, which is why the test suite can cover it without a browser.

06 What a developer would pick up

The prototype is complete as a prototype. Turning it into a hosted service is a well-bounded piece of work, and the seams are already cut.

Work item	What it involves	Effort ESTIMATE
Database-backed storage	Implement <code>CaseStore</code> and the user store against PostgreSQL (about a dozen methods each). Nothing above that interface changes.	5–8 days
EU hosting, TLS, CI/CD	Frankfurt region; container or managed platform; automated deploys; staging environment.	3–5 days
Document storage	S3-compatible EU object storage with encryption at rest, replacing local folders; retention/deletion jobs.	2–3 days
Payments	Stripe or Mollie for the €390 tier and advisory invoices; order state already exists in the product and is currently recorded as "to be invoiced".	3–5 days
Transactional email	Registration, password reset, tax-advisor invitations, report-ready notifications.	2–3 days
Claude reading in production	The integration is written and tested against a stand-in, but has never called the live API . Needs real-document testing, an evaluation set of real layouts, cost and latency measurement, retry handling.	5–8 days
Security hardening	Two-factor login, secrets management, centralised logging with alerting, backups with a restore test, penetration test.	4–6 days
Polish	Mobile refinement, accessibility pass, error states, analytics on the funnel.	3–5 days

Total: roughly 27–43 developer-days — about six to nine weeks for one competent full-stack developer, excluding the legal work that runs in parallel (\$7.2).

A developer joining this reads three files first: `workflow.py` (every operation the product performs), `casefile.py` (the storage interface to replace) and `webapp/server.py` (the HTTP layer). The domain logic in `scorecard.py` and `remediation.py` is where the intellectual property lives and should be changed only with the domain owner — those files carry regulatory consequences, and the tests enforce them.

07 Challenges, stated honestly

7.1 Technical

Challenge	Detail and mitigation
Reading real annual accounts	The hardest technical problem in the product. 33 fields are extracted correctly from four fictional companies, but real German annual accounts vary enormously in layout, and scans are worse. <i>Mitigation, already built:</i> deterministic consistency checks, mandatory human confirmation of every figure, and manual entry always available — a misread is caught, not propagated.
The Claude path is untested live	The request shape is pinned by tests using a stand-in, but no real document has ever been sent. Cost, latency and accuracy on real layouts are unknown.
DATEV mapping uncalibrated	The SKR04 account-range mapping needs validating against three real tax-advisor exports. SKR03 deliberately raises an error rather than guessing.
Six of ten scorecard factors are convention	Equity ratio, EBIT margin and quick ratio are anchored to the Bundesbank distribution of German SMEs. Debt service capacity, leverage, interest cover, overdraft use, reporting cadence and payment behaviour have no published series and can only be calibrated on real placement outcomes.
Thresholds are convention, not calibration	Scorecard breakpoints follow published lending convention. Only the outcome log can validate them, and that takes real engagements.
Single-developer product	~9218 lines written by one person. Well tested and documented, but bus-factor one until someone else reads it.

7.2 Regulatory and legal THE CRITICAL PATH

Area	Position	Open item
§32 KWG (banking licence)	AVOIDED BY DESIGN	No lending, no deposits, never touches client capital.
EU CRA Regulation	AVOIDED BY DESIGN	Transparent rules, no ML score, no rating vocabulary, disclaimer on every report — enforced by automated tests. The free check's self-serve wording still needs a lawyer's read.
EU AI Act	Designed outside Annex III 5(b)	Confirm in writing, in particular for sole proprietors, who are natural persons.
§34c GewO	Decision deferred	Only triggered by a success fee. Paid legal opinion before that fee structure is offered.
GDPR	Product controls in place	Processor agreements, transfer basis for the AI provider, DPIA, records of processing, retention schedule, lawyer-reviewed consent texts.
Professional liability	Not yet in force	Vermögensschadenhaftpflicht before the first paying client.

7.3 Commercial — the risks that actually decide this

- **The central assumption is unvalidated.** That one-third to one-half of obstacle cases are fixable rather than genuinely uncreditworthy is judgement, not data. If the true share is small, the addressable market shrinks by an order of magnitude. This is the first thing the pilot measures, and the outcome log is designed to measure it case by case.
- **Conversion from free to €390 is unknown.** The free check is an acquisition bet. If the band and three weaknesses are enough for the owner, they may never buy the report. The counter-design: the free result deliberately names the weaknesses without quantifying what fixing them is worth.
- **Acquisition channel unproven.** Tax advisors are the natural referral route — they see the rejection first — but they may see an advisory product as competition rather than a service to their clients.
- **Trust.** An SME handing its complete financials to an unknown brand is a real barrier, which is why the security page states plainly what is and is not yet in place rather than making claims.
- **Founder capacity.** Phase 0 is solo and side-structured. Every paid engagement consumes founder days, which is exactly why the free tier had to be self-serve.

- **Timing.** The product is counter-cyclical — demand rises as lending tightens — but if credit conditions loosen materially, urgency falls.

08 Costs

All figures below are planning estimates in EUR, to be verified against current provider pricing. **The infrastructure is not the constraint; legal work and customer acquisition are.**

8.1 One-off, to reach a hosted service with a first paying client

Item	Estimate	Note
Engineering (27–43 days, \$6)	€0 — or €16k–39k	€0 if founder-built; €600–900/day if contracted out
Data-protection package: privacy notice, consent texts, DPIA, records of processing	€2,000–5,000	Specialist lawyer or external DPO
§34c GewO opinion + terms of business	€1,000–2,000	Only needed if a success fee is offered
Professional liability insurance (first year)	€800–2,000	Annual, scales with cover
Penetration test	€2,000–5,000	Can be deferred past the first pilot clients
Bundesbank benchmark data work — done	€0	Seven editions imported; the publication is free
Brand, domain, company formation	€500–1,500	Depends on legal form
Total to a defensible launch	€6k–16k founder-built €22k–55k fully contracted	Excluding founder time

8.2 Running cost, per month at pilot volume

Item	Estimate	Note
EU application hosting + managed PostgreSQL	€40–120	Frankfurt region
Object storage, backups	€10–30	Documents are small
Domain, business email, error monitoring	€15–55	
AI reading and explanations	€0.10–0.60 per check	One document read plus explanations. At 100 free checks a month: €10–60 . Verify against current API pricing; the local reader costs nothing
Payment processing	~1.5% + €0.25	Per transaction, EU cards — about €6 on a €390 report
Fixed monthly run rate	€65–205	Plus usage. One €390 report a month covers it several times over

8.3 Unit economics ILLUSTRATIVE

Product	Price	Marginal cost	Human time
Free check	€0	~€0.10–0.60	None — fully self-serve
Full report	€390	~€7 incl. payment fee	30–45 min advisor review before release
Advisory	€1,500–3,000	negligible	1–2 days per engagement

The shape that matters: the free tier is cheap enough to run at volume without funding it, the report is near-pure margin once the advisor review is efficient, and advisory is capacity-limited rather than cost-limited. The binding constraint at every stage is advisor hours, which is why the first hire in Phase 1 is a junior credit analyst rather than a salesperson.

09 What we want to achieve

The market we are aiming at

Stage	Estimate	Basis
Small and medium enterprises, 10–249 employees (Germany)	~419,000	Destatis / Statista, 2022
Apply for a bank loan in a survey period	~113,000	ECB SAFE (27%) — the free check's market
Face significant obstacles — rejected, partly rejected, discouraged	~16,000	ECB SAFE (14% of applicants)
Fixable rather than genuinely uncreditworthy	~5,000–8,000/yr	UNVALIDATED ASSUMPTION

Stated plainly: euro-area rates applied to Germany (likely an understatement, given Germany's faster deterioration); rolling survey periods, so this is an order-of-magnitude snapshot; the fixable share is judgement and is the first thing to validate.

Milestones

Horizon	Goal
0–3 months	10 pilot cases run free on real company data. Conversations with 5 tax-advisory firms. Legal package commissioned. DATEV mapping calibrated on real exports.
3–6 months	Hosted in the EU with payments live. First paid reports sold. Free-to-paid conversion measured for the first time. Claude reading validated on real documents or dropped in favour of the local reader.
6–12 months	Two to three paid engagements a month. Outcome log at 30+ rows. First recalibration of scorecard thresholds against real placements. First working relationships with specific lenders. Decision on the §34c success-fee model.
Beyond	Junior credit analyst hired for volume; founder on clients and lender relationships; Austria as the second market; the outcome dataset as the defensible asset.

The metrics that decide whether this works

Fixable share Of cases that hit obstacles, how many are presentation rather than substance? The whole thesis.	Free → paid conversion What share of completed free checks buy the €390 report?	Band uplift Measured before/after improvement on real files, not simulated.	Placement rate Of prepared files taken to a lender, how many are approved?
---	---	---	--

10 Status today

BUILT AND TESTED

- Diagnostic engine: 26-value ratio set, 10-factor transparent scorecard, 10 fixability rules, before/after simulation, 7 lender types
- Free AI-assisted credit check: reading, human confirmation, instant result, guarded explanation
- Automatic reading of DATEV trial balances and bank-statement exports
- Public website, investor page and security page, DE/EN
- Three-party portal with server-enforced permissions, guided intake, uploads, letters, advisor release
- Outcome log from engagement one; AI-use audit log; self-service data deletion
- Sector benchmarks and three scorecard factors calibrated against the Bundesbank Jahresabschlussstatistik (firm-level quartiles, by sector and revenue class)
- 217 automated tests, including regulatory-wording, real-HTTP permission and calibration-drift tests
- Verified end to end: a case assembled from uploaded questionnaires, DATEV exports and a bank statement reproduces the hand-built reference case exactly

HONESTLY STILL OPEN

- No paying clients — validated on fictional cases only
- The fixable share (5,000–8,000/yr) is unvalidated
- DATEV mapping uncalibrated; six of ten scorecard factors still rest on convention rather than data
- Claude reading never run against a real document

- No database, no hosting, no payment provider, no transactional email
- Data processing agreements, DPIA, retention schedule, EU hosting, encryption at rest
- §34c GewO decision; EU AI Act classification confirmed in writing; liability insurance

What this is, precisely: a complete, working, tested product that has never met a real customer or a real lawyer. Everything that would be expensive to change later — the regulatory posture, the separation of AI from assessment, the refusal to simulate away genuine risk, the outcome log — is already decided and built. Everything still open is either a purchase, a contract, or a well-bounded engineering task.

11 Appendix — running it locally

For the developer. No dependencies to install, no build step, no services to start.

```
# Demo: website, investor page and portal with four fictional companies
python -m credit_readiness serve --demo --open          # http://127.0.0.1:8766

# Tests
pip install -e ".[dev]" && pytest                      # 204 passing

# A complete engagement on generated fictional data
python scripts/generate_intake_samples.py              # questionnaires, DATEV, bank CSV, PDFs
python scripts/demo_intake_flow.py                    # case → uploads → letters → diagnostic → outcome
```

Demo login (password <code>demo1234</code>)	Role	Shows
<code>jonas.weber@demo.de</code>	Company	The free credit check from step one — upload an annual-accounts PDF and watch it work
<code>anna.mueller@demo.de</code>	Company	A complete engagement, report released: band C → A, fixable
<code>elif.yilmaz@demo.de</code>	Company	The honest no: genuine credit risk, advised not to apply
<code>kanzlei@demo.de</code>	Tax advisor	Only its own clients' cases; confirmations and sign-off sheet
<code>berater@demo.de</code>	Advisor	The pipeline: run analyses, release reports, record outcomes

Delete `data/demo/` to reset the demo; it is re-created on the next start. Real use runs against `data/clients/`, which is excluded from version control. Sample PDFs for testing the reader are in `data/samples/intake/`.

Disclaimers. This document describes a pre-revenue product. Market figures come from the sources named; the fixable share, all fee levels, all cost estimates and all conversion assumptions are unvalidated planning assumptions. All demo companies and all sample documents are fictional and must never be presented as evidence about the real market. The product produces a directional readiness indication — never a credit rating, a probability of default, or a promise of approval. The lender decides. Nothing here is legal advice; the items marked open in §7.2 require a qualified lawyer before real client data is processed on a server.